

AI Capability Practice Guide: Individual Practice

Practical, Responsible, High-Impact AI Use for Everyday Professional Work

Who This Guide Is For

This guide is for professionals, practitioners, and knowledge workers who are already using—or are about to use—AI tools in their everyday work and want to do so:

- confidently
- responsibly
- effectively
- without exposing themselves or others to unnecessary risk

You do **not** need technical expertise.

You do **not** need to study AI theory beforehand.

You **do** need to make sound decisions under real-world pressure.

This guide assumes you are the person responsible for the outcome—even if AI is involved.

Who This Guide Is Not For

This guide is not designed for:

- technical model development or AI engineering
- performance tuning or optimisation tutorials
- “growth hacking” or automation-at-all-costs approaches
- replacing professional judgement with AI output

If you are looking for shortcuts that remove responsibility, this guide will feel deliberately uncomfortable.

What You Will Be Able to Do in 30–60 Minutes

By working through this guide, you will be able to:

- decide *when* AI use is appropriate—and when it is not
- design a clear human–AI working relationship for a real task
- identify ethical, equity, and impact risks before harm occurs
- apply lightweight governance to your own AI use
- document decisions in a way you can justify later
- reflect on outcomes and improve your practice over time

You will also produce at least one **tangible artefact** (a documented decision, revised workflow, or reflection record) you can reuse immediately.

FAST START — USE THIS NOW

If you only read one section, read this one.

This Fast Start is designed to let you **use AI responsibly within 10 minutes**—without reading the rest of the guide first.

When to Use This Guide

Use this guide when:

- you feel pressure to “just use AI” because others are doing so
- a task matters, but the rules are unclear
- you are not sure how much to trust an AI output
- the impact extends beyond you (colleagues, clients, students, the public)
- you might be asked later to justify how a decision was made
- you sense risk, but can’t quite articulate where it lies

If none of these are true, this guide may not be necessary yet.

The 10-Minute Entry Workflow

Use this sequence *before* you open an AI tool or act on its output.

Step 1 — Name the task

Write down, in plain language:

“I am using AI to help me with: [task]”

Avoid abstract phrasing. Be concrete.

Step 2 — Decide the role of AI (Co-Agency check)

Ask:

- What part of this task can AI support?
- What part **must remain human-led?**

If you cannot clearly answer both, stop and refine the task before proceeding.

Step 3 — Apply the primary capability check

Ask:

- Do I understand how this AI might produce an answer?
- Do I know what it *might get wrong* here?

If not, you are not ready to rely on the output.

Step 4 — Run the rapid ethics & impact screen

Ask:

- Who could be affected by this output besides me?
- Could this introduce bias, misrepresentation, or harm?

If impact extends beyond yourself, increase scrutiny.

Step 5 — Decide the action

Choose **one**:

- Proceed with AI support
- Revise the task or prompt
- Pause and escalate (seek review, expertise, or decide not to use AI)

Document the choice. Even one sentence is enough.

Worked Example: Same Task, Three Outcomes

Task

Using AI to draft a short professional summary for public use.

Good Use

- AI produces a draft
- Human reviews for accuracy, tone, and context
- Sensitive claims are verified or removed
- Final wording is human-authored and owned

Why this works:

AI accelerates drafting without replacing judgement or accountability.

Risky Use

- AI output is lightly edited
- No check for assumptions or omissions
- Context-specific nuance is missing

Why this's risky:

Errors aren't obvious, but credibility could be undermined later.

⊖ Unacceptable Use

- AI output is published as-is
- No awareness of audience impact
- No accountability if challenged

Why this fails:

Responsibility has been delegated to a system that cannot carry it.

Your First Artefact (Create This Now)

Write a short **AI Use Note** (3–5 lines):

Task:

Role of AI:

Human responsibility:

Risks considered:

Decision made:

This single note is enough to:

- clarify your own thinking
- demonstrate responsible practice
- protect you if questions arise later

You have now already improved your AI capability.

HOW THIS GUIDE WORKS

This guide is designed to be **used**, not read from start to finish.

You are not expected to move linearly. You are expected to:

- dip in when a situation arises
- apply a tool
- make a judgement
- move on

This section explains **how to navigate the guide without increasing cognitive load**.

The Six Domains as a Living Workflow

The AI Capability Framework is built around six domains.

In this guide, they operate as a **practical workflow**, not a conceptual model.

You do **not** apply all six domains equally every time.

You apply them **as the situation demands**.

The capability flow in practice

- **AI Awareness & Orientation**
Understand what the system is doing and where it might mislead.
- **Human–AI Co-Agency**
Decide who does what, and who remains accountable.
- **Applied Practice & Innovation**
Use AI creatively and productively, but within boundaries.
- **Ethics, Equity & Impact**
Anticipate who may be affected and how harm could occur.
- **Decision-Making & Governance**
Apply proportionate oversight, documentation, and review.

- **Reflection, Learning & Renewal**

Learn from outcomes and improve your future practice.

Skipping a domain does not save time.

It usually **moves risk downstream**, where it is harder to see and harder to correct.

How to Use This Guide Under Time Pressure

Most AI decisions are made:

- mid-task
- under delivery pressure
- with incomplete information

This guide is structured to support exactly that reality.

If you only have 5–10 minutes

- Use the **Fast Start**
- Apply the co-agency check
- Run the rapid ethics screen
- Document the decision briefly

That alone significantly improves your practice.

If you have 20–30 minutes

- Identify which **situational entry point** fits your context
- Apply the relevant domains
- Use one template or checklist

- Capture one reflection insight

This is the *recommended minimum* for non-trivial work.

If stakes are high

- Work through all domains
- Pay particular attention to ethics and governance
- Use the escalation guidance later in the guide
- Document reasoning clearly

High-impact contexts require *slower, more deliberate capability*, not faster automation.

SITUATIONAL ENTRY POINTS

Start where your problem is — not where the framework begins.

Most people don't begin by asking,
“Which AI capability domain should I apply?”

They begin with a **situation**.

Use the entry point below that best matches your reality. Each points you to what to do next.

Entry Point 1 — “I need to move quickly, but I’m uneasy.”

You feel pressure to deliver fast, but something about using AI feels risky.

Primary domains to apply

- Human–AI Co-Agency
- Ethics, Equity & Impact

What to do now

- Clarify which parts of the task must remain human-led
- Identify who could be affected if the output is wrong
- Increase review before sharing or acting on the output

Common failure mode

- Speed displacing judgement
 - Assumption that quick work is low-impact
-

Entry Point 2 — “I’m not sure if AI is appropriate here.”

You could use AI — but should you?

Primary domains to apply

- AI Awareness & Orientation
- Decision-Making & Governance

What to do now

- Identify what kind of system you are interacting with
- Ask what it cannot reasonably be expected to know
- Decide whether AI should assist, inform, or be excluded

Common failure mode

- Treating uncertainty as a signal to “try anyway”
-

Entry Point 3 — “The output looks good, but I don’t fully trust it.”

The AI result appears polished, but confidence is low.

Primary domains to apply

- AI Awareness & Orientation
- Reflection, Learning & Renewal

What to do now

- Identify assumptions or gaps in the output
- Cross-check critical claims or framings
- Adjust your future prompts or workflows

Common failure mode

- Mistaking fluency for reliability
-

Entry Point 4 — “This could affect other people.”

The work goes beyond personal use.

Primary domains to apply

- Ethics, Equity & Impact
- Decision-Making & Governance

What to do now

- Identify impacted groups
- Consider fairness, bias, and misrepresentation
- Decide what oversight or disclosure is appropriate

Common failure mode

- Treating indirect impact as negligible
-

Entry Point 5 — “Someone might question how this was produced.”

You anticipate scrutiny — now or later.

Primary domains to apply

- Decision-Making & Governance
- Human–AI Co-Agency

What to do now

- Document how AI was used and where humans intervened
- Ensure accountability is clear
- Avoid delegating judgement retroactively

Common failure mode

- Trying to reconstruct decisions after the fact
-

Entry Point 6 — “I want to get better, not just faster.”

You’re thinking about long-term capability, not one task.

Primary domains to apply

- Reflection, Learning & Renewal
- Applied Practice & Innovation

What to do now

- Identify patterns in what works and what fails
- Adjust how you design tasks and prompts
- Capture insights for future reuse

Common failure mode

- Repeating mistakes without noticing them
-

How the Rest of This Guide Is Structured

From this point onward, the guide moves into the **core practice workflow**, where each domain is treated as a practical instrument.

Each domain section will:

- explain what it protects or enables
- show how to apply it immediately
- identify common failure modes
- include a short reflection moment

You can engage with:

- one domain
- several domains
- or all six

depending on the situation you face.

CORE PRACTICE WORKFLOW

Domains 1–3: Working Well With AI

These first three domains shape *how you engage with AI at the point of use*. They focus on understanding, role clarity, and responsible application.

Domain 1 — AI Awareness & Orientation

What This Domain Protects

This domain protects you from:

- false confidence in AI outputs
- misunderstanding what an AI system can or cannot know
- treating generated content as authoritative when it is not

AI systems do not “know” in human terms.

They generate outputs based on patterns in data and probability.

Awareness is what allows you to decide **how much trust is appropriate**.

Apply Now — Key Questions

Before using or acting on an AI output, ask:

- What kind of system am I interacting with?
- What information *might be missing* from its training or context?
- What assumptions is the output making implicitly?
- Where would an error matter most?

If you cannot identify plausible failure points, you are over-trusting the output.

Tool in Use — Awareness Check

Use this quick check before relying on an output:

AI Awareness Check

- What is the AI optimised to do here?
- What is it *not* designed to do?
- What would a confident-sounding but wrong answer look like?

Capture one sentence for each.

Common Failure Modes

- Confusing fluency with accuracy
 - Assuming recent or niche knowledge without verification
 - Treating the output as neutral or objective
-

Quick Reflection

What did I assume the AI “knew” that it could not reasonably know?

Write one sentence. This improves future judgement immediately.

Domain 2 — Human–AI Co-Agency

What This Domain Protects

This domain protects:

- accountability
- professional responsibility
- ethical ownership of decisions

AI can assist, but it cannot carry responsibility.
Co-agency makes that boundary explicit.

Apply Now — Key Questions

Ask:

- Which parts of this task can be delegated to AI support?
- Which parts **must remain human-led**?
- Who is accountable if something goes wrong?

If accountability is unclear, co-agency is poorly designed.

Tool in Use — Co-Agency Map

Define roles explicitly:

Co-Agency Map

- AI supports by:
- Human decides on:
- Human remains accountable for:

This takes under two minutes and prevents role drift.

Common Failure Modes

- Letting AI shape decisions implicitly
 - Accepting suggestions without scrutiny
 - “Rubber-stamping” AI outputs
-

Quick Reflection

Did I actively design the relationship—or let it emerge by default?

Domain 3 — Applied Practice & Innovation

What This Domain Enables

This domain enables:

- productive experimentation
- creative use of AI
- improvement of workflows

Innovation here means **better practice**, not maximal automation.

Apply Now — Key Questions

Ask:

- Does AI actually improve this task?
- What would “good” look like without AI?
- Am I experimenting safely, or exposing others to risk?

Not every task benefits from AI assistance.

Tool in Use — Safe Experimentation Prompt

When exploring AI use, try:

“Generate a draft to support ideation. I will review for accuracy, bias, and appropriateness before use.”

Explicitly frame AI as *support*, not authority.

Common Failure Modes

- Using AI because it is available, not because it adds value
 - Scaling use before understanding consequences
 - Treating experimentation as risk-free
-

Quick Reflection

What improved because I used AI—and what didn’t?

Capture one insight for reuse later.

Domains 1–3 in Practice

Together, these domains shape:

- how you understand AI output
- how you maintain control and responsibility
- how you innovate without carelessness

They are **necessary but not sufficient** for responsible practice.

As impact and risk increase, you must apply the next domains.

RISK, RESPONSIBILITY & RENEWAL

Domains 4–6: Governing Impact and Sustaining Capability

Domains 4–6 apply when AI use affects other people, shapes decisions, or persists over time.

These domains turn good use into *responsible use*.

Domain 4 — Ethics, Equity & Impact

What This Domain Protects

This domain protects:

- people affected by AI-influenced decisions
- institutional trust
- social legitimacy and fairness

Ethics is not an after-the-fact check.

It is a **design constraint**.

Apply Now — Key Questions

Before sharing, publishing, deploying, or acting on AI output, ask:

- Who benefits from this output?
- Who might be disadvantaged, misrepresented, or excluded?
- What assumptions about people, culture, or context are embedded?
- What harm would be hardest to reverse?

If potential harm exists and is unexamined, pause.

Tool in Use — Rapid Ethical Check

Use this 90-second scan:

Ethical Impact Scan

- Accuracy risk: What might be misleading or wrong?
- Bias risk: Whose perspectives are missing or flattened?
- Power risk: Who has less ability to challenge this output?

One sentence per risk is sufficient.

Common Failure Modes

- Treating ethics as compliance
 - Assuming neutrality because the output “sounds reasonable”
 - Overlooking who cannot contest or correct the output
-

Quick Reflection

If I were on the receiving end of this output, what would concern me most?

Domain 5 — Decision-Making & Governance

What This Domain Protects

This domain protects:

- accountability
- institutional legitimacy
- defensibility of decisions

Governance answers one question:

“How was this decision shaped, and who takes responsibility?”

Apply Now — Key Questions

Ask:

- Does AI influence this decision directly or indirectly?
- Is this a low-risk assistive use or a high-impact judgement?
- Who can challenge the outcome?

If you cannot document these answers, governance is missing.

Tool in Use — Decision Transparency Log

For any AI-influenced judgement:

Decision Transparency Log

- Decision being made:
- Role of AI (inform / suggest / generate):
- Human decision-maker:
- Review or escalation required? (Yes / No)

This log can be informal—but must exist.

Risk Thresholds (Practical Rule)

Escalate or require additional review when AI use involves:

- public communication
 - assessment, grading, or evaluation
 - clinical, legal, or safety implications
 - reputational or policy consequences
-

Common Failure Modes

- Letting AI quietly reshape decisions
 - Assuming speed equals improvement
 - No audit trail for future challenge
-

Quick Reflection

Could I clearly explain and justify this decision six months from now?

Domain 6 — Reflection, Learning & Renewal

What This Domain Sustains

This domain sustains:

- long-term capability
- ethical maturity
- adaptive intelligence

Without reflection, capability stagnates—even when tools improve.

Apply Now — Key Questions

After meaningful AI use, ask:

- What worked better because of AI?
- What new risks appeared?
- What should we change next time?

Reflection turns experience into capability.

Tool in Use — Mini Reflection Cycle

Use this short loop:

Reflect → Adjust → Reapply

- One insight
- One improvement
- One boundary reset

This can take under five minutes.

Common Failure Modes

- Repeating the same prompts uncritically
 - Treating mistakes as personal failure instead of learning signals
 - No collective learning across teams
-

Quick Reflection

What assumption about AI shifted for me during this task?

Domains 4–6 in Practice

Together, these domains ensure that:

- innovation does not outrun responsibility
- decisions remain contestable and transparent
- learning compounds rather than resets

They are essential for:

- teaching
 - research
 - leadership
 - public or organisational impact
-

The Full Capability Cycle (Operational View)

AI capability matures through repetition of this loop:

Awareness → Co-Agency → Practice → Ethics → Governance → Reflection → Renewal

Skipping steps increases risk.

Revisiting steps deepens capability.

STAGE 5 — CAPABILITY SELF-CHECK & WORKED EXAMPLE

Making Capability Visible and Actionable (Individual Practice)

Part A — AI Capability Self-Check (Individual)

This is **not an assessment** and **not a scorecard**.

Its purpose is orientation: *Where am I right now, and what should I focus on next?*

Complete this in **under five minutes**.

Domain 1 — Awareness & Orientation

Answer honestly:

- I know how AI generates responses (at a basic level)
- I routinely check accuracy rather than trusting fluency
- I can explain AI limitations to someone else

If mostly NO → start with Domain 1 actions before increasing usage.

Domain 2 — Human–AI Co-Agency

Consider:

- I decide *when* to use AI rather than defaulting to it
- I am clear what I delegate to AI vs what I retain
- I remain accountable for final outputs

If **unclear** → prioritise co-agency boundaries.

Domain 3 — Applied Practice & Innovation

Reflect:

- I use AI to explore options, not just get answers
- I iterate rather than accept first outputs
- I treat AI as a thinking partner, not a shortcut

If **limited** → build safe experimentation habits.

Domain 4 — Ethics, Equity & Impact

Ask yourself:

- I consider who may be affected by AI-influenced outputs
- I can identify potential bias or misrepresentation
- I pause when consequences are hard to reverse

If **rarely** → ethics must move earlier in your workflow.

Domain 5 — Decision-Making & Governance

Check:

- I can explain how AI influenced my decisions
- I know when escalation or review is needed
- I avoid “invisible” AI influence on judgement

If missing → document decisions before risk increases.

Domain 6 — Reflection & Renewal

Finally:

- I review AI use outcomes deliberately
- I update prompts and boundaries over time
- I learn from failures, not just successes

If inconsistent → adopt a lightweight reflection habit.

Interpreting Your Self-Check

- **Most gaps in Domains 1–2:** slow down and build foundations
- **Most gaps in Domains 4–5:** increase oversight before scaling use
- **Most gaps in Domain 6:** capability will stagnate without renewal

Capability grows by *rebalancing*, not maximising.

Part B — Worked Example: Individual AI Practice in Action

This example is deliberately **ordinary**.
No specialist role. No advanced tooling.

Scenario

You must produce a **high-stakes written output** under time pressure:

- a professional report
- a critical briefing
- a proposal or analysis

You decide to use AI.

Domain 1 — Awareness in Action

You begin by recognising:

- The model is probabilistic, not authoritative
- It may fabricate plausible-sounding details
- Fluency ≠ accuracy

You therefore:

- avoid asking for “final answers”
 - ask for structured drafts and alternatives
-

Domain 2 — Co-Agency in Action

You define roles clearly:

AI does:

- suggest structure
- surface alternative framings
- help draft early versions

You do:

- validate claims
- select arguments
- own conclusions

This prevents silent delegation.

Domain 3 — Applied Practice in Action

You use AI to:

- compare multiple ways to frame the issue
- test clarity with different audiences
- generate counter-arguments

You **do not**:

- copy outputs uncritically
- bypass your own reasoning

Iteration replaces extraction.

Domain 4 — Ethics & Impact in Action

Before sharing the output, you ask:

- Could this mislead or oversimplify?
- Are key perspectives missing?
- Who might be disadvantaged by how this is framed?

You adjust tone, emphasis, and qualifiers accordingly.

Domain 5 — Governance in Action

You ensure:

- AI use could be declared if required
- You can explain how AI influenced the work
- High-risk claims were reviewed manually

If challenged later, the decision trail exists.

Domain 6 — Reflection in Action

After completion, you reflect:

- What did AI genuinely improve?
- Where did it introduce risk or noise?
- What will I change next time?

You update:

- prompts
- boundaries
- decision thresholds

Capability compounds.

What This Example Shows

AI capability is not about *what tool you used*.
It is about **how you designed the interaction**.

The same logic applies across:

- teaching
- research
- leadership
- public impact

This is the **individual foundation** all other contexts build on.

STAGE 6 — QUICK-START & PERSONAL OPERATING MODEL

The Individual AI Capability Playbook (At-a-Glance)

Part A — One-Page Quick-Start (Immediate Use)

If you only read one page of this guide, read this.

The Individual AI Capability Loop

Use this loop **every time AI matters**:

Awareness → Co-Agency → Practice → Ethics → Governance → Reflection → Renewal

Skipping steps increases risk.

Repeating steps increases capability.

The 60-Second Pre-Use Check

Before using AI, ask:

1. **Why am I using AI here?**
(Speed, clarity, exploration — not avoidance.)
2. **What will AI do — and what remains mine?**
(Define the boundary.)
3. **What would failure look like?**
(Accuracy, harm, reputation, fairness.)

If you cannot answer these, pause.

The 3 Rules of Individual AI Practice

1. **Fluency is never proof.**
Always verify claims that matter.
 2. **You remain accountable.**
AI assistance does not dilute responsibility.
 3. **If impact extends beyond you, governance applies.**
Document or escalate.
-

Common Signals to Slow Down

- Outputs sound “too confident”
- Decisions feel easier but less examined
- You could not explain the reasoning later

These are **capability warnings**, not tool failures.

Part B — Personal AI Practice Operating Model

This is your **repeatable system**, not a checklist.

Complete this once, then revisit quarterly.

1 My Valid Reasons to Use AI

Examples:

- clarifying complex material
- exploring alternative framings
- drafting early versions
- sense-checking logic

I do not use AI to avoid thinking or responsibility.

Write yours here:

2 My Co-Agency Rules

AI may:

AI may not:

I always retain responsibility for:

This prevents unconscious delegation.

3 My Ethical Red Lines

I pause or stop AI use when:

- _____
- _____

Examples:

- risk of misrepresentation
 - impact on vulnerable groups
 - irreversible public consequences
-

4 My Governance Triggers

I document or escalate AI use when:

- _____
- _____

Examples:

- assessments or evaluations
 - public communication
 - policy, safety, or reputational risk
-

5 My Reflection Habit

After meaningful AI use, I ask:

- What helped?
- What harmed or confused?
- What will I adjust next time?

My reflection cadence:

- after every task
 weekly
 project-based
-

6 My Renewal Commitments

To keep capability current, I will:

- revisit assumptions quarterly
- update prompts and boundaries
- stay alert to new risks

Capability decays without renewal.

The Individual Promise

I use AI deliberately, not reflexively.

I design my collaboration with AI.

I remain responsible for outcomes.

I reflect so my capability grows over time.
